

The Agentic Revolution: A Report on Modern Agentic Coding Practices and Documentation Standards

1. Executive Summary: The Agentic Revolution in Software Development

1.1. Introduction to the Paradigm Shift

The landscape of software development is undergoing a fundamental transformation, moving beyond traditional, human-centric coding practices toward a new paradigm of agentic orchestration. This shift is characterized by the emergence of autonomous, goal-oriented AI systems—known as agents—that are capable of performing multi-step tasks with minimal human intervention.¹ These systems are far more sophisticated than previous generations of AI tools, which were primarily limited to reactive functions like code completion. The move to agentic coding redefines the developer's role from writing individual lines of code to designing and managing entire AI-driven systems.

This report synthesizes the latest practices in agentic development, identifying core principles, architectural standards, and key frameworks that are shaping this new era. It specifically examines the rise of AGENTS.md as a critical standard for machine-readable project instructions and analyzes the paramount importance of stateful memory for creating truly autonomous, self-improving systems. The analysis concludes with a comparative overview of leading agentic frameworks, including a focused case study on Roo Code, to provide a blueprint for documentation and implementation.

1.2. Key Findings and Recommendations

- **AGENTS.md as the De Facto Standard:** The analysis confirms AGENTS.md is an increasingly adopted open standard for providing structured, machine-readable instructions to coding agents. Its adoption by major players like Google, OpenAI, and Sourcegraph points to its role in fostering a multi-vendor, interoperable ecosystem. The report recommends its use for all new projects to ensure consistency, improve code quality, and accelerate agent onboarding.
- **The Primacy of Stateful Memory:** The report establishes that effective long-term memory is the key to advancing beyond stateless, single-session interactions. By enabling agents to learn from past experiences and maintain knowledge across sessions, stateful architectures provide the necessary foundation for true agentic autonomy and self-evolution.
- **Architectural Nuance over Monolithic Tools:** The evidence suggests that best practices are not defined by a single "best" framework but by a developer's ability to choose the right architectural philosophy for a given task. This involves understanding the trade-off between predictable, structured workflows and dynamic, autonomous agents. The report recommends a hybrid approach that combines the strengths of both.
- **Roo Code as a Privacy-Focused Model:** The report identifies Roo Code as a prime example of a broader trend toward local, transparent, and privacy-centric agentic tools. Its open-source, local-first architecture directly addresses the critical enterprise need for control over proprietary and sensitive code, making it a viable alternative to black-box, cloud-based solutions.

2. The Agentic Coding Paradigm: A Shift to Orchestration

2.1. Defining Agentic Coding and Its Core Principles

Agentic coding represents a fundamental shift in the relationship between humans and AI in the software development process. It moves beyond the reactive, autocomplete-style assistance of early AI tools toward a proactive, multi-step, and goal-oriented approach. At its core, agentic coding involves creating systems where AI agents, which are autonomous and goal-oriented programs powered by large language models (LLMs), can take on complex

tasks with minimal human intervention.¹ The overarching objective is to free developers from the intricate details of the "how" of coding, allowing them to focus on the strategic "why" and the architectural "what".¹

True agentic systems are built on four key principles:

- **Autonomy:** The agent possesses the capacity to make its own decisions on how to achieve a given goal. A developer provides the destination, and the agent independently plots the route, adapting its course as needed.¹ This ability to operate without constant human supervision is a hallmark of the agentic paradigm.⁴
- **Tool Use:** Agents are not confined to merely generating text or code. They are equipped with the ability to interact with the digital world by executing code, calling APIs, searching the web, and reading or writing files. This tool-use capability is what transforms a language model into an operational agent.¹
- **Multi-Step Reasoning:** Complex problems are not solved in a single step. Agentic systems excel at breaking down an ambiguous, large problem into a sequence of smaller, manageable sub-tasks. They follow a logical "chain of thought," where the output and results from one step inform and guide the next.¹
- **Self-Correction:** A crucial aspect of autonomous operation is the ability to handle errors. An agent can test its own output, recognize a problem, and loop back to try an alternative approach. This self-debugging capability allows it to refine its work without requiring human intervention for every misstep.¹

2.2. The Evolving Role of the Developer

The adoption of agentic systems necessitates a significant evolution of the developer's role. The primary function shifts from being a hands-on coder who meticulously writes individual parts of a system to becoming a high-level architect who designs and orchestrates the entire system.¹ The developer is no longer the expert who knows everything but rather a perpetual learner and a master collaborator who designs brilliant systems.¹

This new role emphasizes several key skills:

- **Systems Thinking:** The developer's primary job is to deconstruct a large, ambiguous goal into a logical workflow of discrete tasks. This involves architecting the "orchestra" of agents, determining which agent performs which task, in what order, and how data is handed off between them. The developer must also design robust mechanisms for recovering from potential failures.¹
- **High-Level Testing:** When an AI agent can generate a thousand lines of code, reviewing it line by line becomes an untenable practice.¹ As a result, the developer must shift their

testing strategies. The best practice is to create "super-tests," which are high-level validation scripts that ensure the agent's output meets the functional and non-functional requirements of the system as a whole. This moves the focus from granular unit testing to comprehensive system-level validation and adversarial testing designed to uncover edge cases.¹

This shift to orchestration underscores a new cognitive skill set for developers. If an agent can handle the intricate details of the "how," the developer's expertise must move higher up the stack. A developer's value is no longer measured by their ability to write the most elegant code but by their proficiency in system architecture, workflow design, and high-level testing. This redefines "full-stack development" from a single person working at all layers of the technical stack to a professional who can design and manage a complex system where different layers are handled by a combination of human and AI agents. The core value of agentic coding is not the generation of code itself but the ability to solve problems at a higher level of abstraction.¹ The ability to design the orchestration logic and the super-tests allows the developer to offload the entire execution of a complex, pre-defined workflow, which is the causal link between the agent's technical capabilities and the evolution of the developer's role.

3. The AGENTS.md Standard: A New Language for AI Collaboration

3.1. Purpose and Value of AGENTS.md

AGENTS.md is an emerging, open standard that is rapidly gaining traction across the software development community. It is a dedicated Markdown file that provides clear, structured instructions for AI coding agents and serves as a predictable, machine-readable complement to the human-facing README.md.⁵ Its simplicity and lack of a complex schema are strategic design choices that have positioned it as a key enabler for a multi-vendor, interoperable ecosystem.⁵

The value of adopting AGENTS.md is multifold:

- **Improved Code Quality and Consistency:** By documenting project conventions—such as coding style, formatting, and architectural patterns—in a structured format, AGENTS.md ensures that AI-generated code conforms to team standards.⁵ This steers

the AI away from bad practices and common pitfalls, reducing bugs and the need for extensive refactoring. For example, if a project requires the use of a specific logging library,

AGENTS.md can explicitly state this, ensuring the AI agent uses the correct module instead of `print()` statements.⁹

- **Accelerated AI Onboarding:** Much like a new human developer, an AI agent needs to learn a project's context and expected behaviors. An AGENTS.md file condenses this knowledge into a format that the agent can parse in seconds, allowing it to become productive in a fraction of the time it would take to read and understand a codebase from scratch.⁹ This is a direct solution to a major productivity bottleneck.
- **Reduced Development Time:** When an AI agent has access to clear, project-specific instructions, its initial output is much more likely to be on-target and compliant.⁹ This means developers spend less time fixing stylistic or architectural issues, which speeds up coding tasks and allows them to focus on more complex, higher-level problems.

The simplicity of AGENTS.md is a powerful catalyst for its broad adoption. It is a simple text-based convention with no special tooling or schema required, which prevents it from becoming a proprietary format that locks developers into a specific vendor's ecosystem.⁵ The community's convergence on this single, open standard provides an elegant solution to the challenge of managing multiple instructions files for different coding assistants.⁷

3.2. Implementing AGENTS.md: Structure and Best Practices

A well-structured AGENTS.md file should be treated as a form of source code and kept up to date alongside the codebase.¹⁰ While there is no strict schema, a robust file typically includes the following sections:

Section Name	Purpose (for the AI agent)	Example Content
Project Overview & Structure	Provides a high-level summary and explains the project's directory layout to help the agent understand where different types of code belong.	The project uses an MVC architecture. Frontend code is in <code>/webapp</code> , the API is in <code>/server</code> , and data models are in <code>/models</code> .
Dev Environment Tips	Lists explicit commands and shortcuts for setting	Use <code>pnpm install</code> to install dependencies. To run a

	up the environment, navigating the codebase, and running common tasks.	package, use <code>pnpm turbo run start --filter=<project_name></code> .
Build and Test Commands	Gives the agent the precise commands needed to build the project and run the full test suite.	Run <code>pnpm turbo run test</code> from the root directory to execute all tests. To run tests for a single package, use <code>pnpm turbo run test --filter=<package_name></code> .
Code Style & Formatting Rules	Details project-specific conventions for naming, formatting, and design patterns.	Use <code>snake_case</code> for all Python function and variable names. All modules must include a docstring at the top.
Security and Dependency Policies	Outlines guidelines for managing dependencies and security considerations.	Always use Python's built-in <code>logging</code> library for debug and error messages instead of <code>print()</code> . Do not install packages from untrusted sources.

For large projects or monorepos, a best practice is to use a modular architecture with nested AGENTS.md files.⁵ An agent will automatically read the file nearest to the code it is working on, which ensures subprojects receive tailored guidance.⁵ This approach prevents the need for a single, giant file with complex conditionals and allows packages with different stacks or versions to evolve independently.⁸

Other key best practices for effective implementation include:

- **Be Concise and Concrete:** Use explicit, bulleted lists and wrap shell commands in backticks.¹⁰ Agents process information more effectively when it is simple and direct, so avoid verbose narratives.
- **Link Rather Than Duplicate:** If detailed information already exists in a README.md or a wiki, simply reference it.¹⁰ This avoids documentation conflicts and reduces the need to maintain duplicate information.

4. Best Practices for Agentic Memory and Persistent State

4.1. The Fundamental Challenge: Beyond the Context Window

A critical limitation of large language models is their inherent statelessness. They possess a vast store of knowledge within their weights but cannot form new memories or learn from experience during deployment.¹² The only information they can recall during an interaction is what fits within their "context window," which functions as a form of ephemeral, short-term memory.¹³ When this window reaches its limit, the model "forgets" the earliest tokens to make room for new ones, leading to a loss of critical context that can cause misinterpretations or factual inaccuracies.¹³

This limitation presents a significant challenge for creating long-running, autonomous agents. A community discussion highlights this issue as the "harder half" of agentic development: managing durable, queryable, and scoped "project memory" that includes decisions, constraints, and deprecations.¹⁵ Without a robust system to manage this persistent knowledge, an agent is limited to a single-session interaction and can confidently cite stale truths.¹⁵

4.2. Architectures for Persistent State

To overcome the limitations of the context window, agentic systems must be designed with architectures that support persistent state. This involves a tiered approach to memory management¹⁴:

- **Ephemeral Session Memory:** This is the most basic form of state, where information is stored only for the duration of a single session. This is sufficient for simple applications like a shopping assistant that needs to track real-time interactions but does not need to persist data for future sessions.¹⁴
- **Persistent Conversation History:** To maintain context across multiple sessions, conversation history can be stored in a durable external database or vector store. This is vital for applications that require long-term memory for personalization, allowing the agent to remember and adapt to a user's preferences over time.¹⁴

- **Structured Knowledge:** For more advanced agents, a common practice is to use vector databases like Qdrant to store semantic knowledge.¹⁶ These databases use embeddings to encode information, enabling the agent to perform semantic searches and retrieve relevant, context-aware information from a vast knowledge base. A common technique is to combine a sliding window for recent messages with periodic summarization of past interactions to control token costs and stay within context limits.¹⁴

4.3. Emerging Best Practice: Agentic Memory (A-MEM)

An emerging best practice moves beyond traditional, developer-defined memory schemas to a more dynamic, agent-driven approach. A recent research paper proposes a novel system called A-MEM, which is a dynamic memory system for LLM agents inspired by the Zettelkasten knowledge management method.¹⁸

This approach is characterized by several core concepts:

- **Dynamic Organization:** Unlike traditional systems that rely on a fixed, predefined structure, A-MEM dynamically organizes memories in an agentic way.¹⁸
- **Note Construction and Linking:** When a new memory is added, an LLM agent automatically generates a structured note with keywords, tags, a contextual description, and links to relevant memories it identifies from its history.¹⁸
- **Memory Evolution:** This is a crucial feature that allows the system to continuously refine its understanding. As new memories are integrated, they can trigger updates to the contextual representations and attributes of existing historical memories, allowing the knowledge network to evolve over time.¹⁸

This dynamic, agent-driven memory system represents a profound shift. The traditional approach requires a human developer to design and manage a fixed database schema. The A-MEM approach, however, delegates the *organization* of memory to the agent itself through a dynamic, Zettelkasten-inspired system. This delegates the responsibility of memory management from the human's "how" (e.g., "design a schema") to the agent's "how" (e.g., "link this memory to existing ones"). This is a microcosm of the larger agentic paradigm shift, where the system itself finds its own solution.¹ This type of stateful architecture is the foundation of true agentic autonomy and self-evolution, as an agent that can learn from its experiences and persist knowledge is no longer a stateless workflow but a continuously improving, long-running system.¹²

5. Comparative Analysis of Leading Agentic Frameworks

5.1. The "Agents vs. Workflows" Dichotomy

A central concept in understanding the agentic framework ecosystem is the distinction between "agents" and "workflows".¹⁹ These two terms are not interchangeable and represent fundamentally different architectural philosophies.

- **Workflows:** These are systems where an LLM is orchestrated through a predefined, deterministic code path.¹⁹ This approach is predictable, reliable, and generally cheaper and faster for tasks that are well-defined and have clear steps. A bug triage system that uses a sequence of agents to analyze a report, assign a priority, and post an update is a classic example of a workflow.¹
- **Agents:** These are systems where the LLM dynamically directs its own process and tool usage, making decisions in real-time about how to accomplish a task.¹⁹ This approach is better suited for flexible, ambiguous problem-solving where the path to a solution is not known in advance.

The evidence suggests that nearly all production agentic systems are a hybrid of both approaches.¹⁹ The best practice is not to choose one or the other but to select a framework based on which of these two philosophies it prioritizes, and then design a system that combines both where appropriate.

5.2. Framework Profiles

The following table provides a high-level comparison of leading agentic frameworks, highlighting their architectural philosophies and intended use cases.

Framework	Architectur al Philosophy	Core Component s	Recommen ded Use Case	Memory Handling	AGENTS.m d Support
-----------	---------------------------------	------------------------	-----------------------------	--------------------	-----------------------

AutoGen	Multi-Agent Conversation & Collaboration	Conversable Agents, User Proxy, Groups	Real-time, scalable systems that require complex, collaborative tasks and multi-agent conversations ²⁰	Short-term memory is built-in; long-term requires external tools like SQLite ¹⁷	Yes, as part of the broader ecosystem of interoperable standards ⁷
CrewAI	Role-Based Orchestration	Crew, Agents, Tasks, Processes	Collaborative teamwork tasks, such as research, content creation, or project management, with a clear division of roles ²⁰	Supports both short-term and long-term memory ¹⁷	Yes, as a key component of enterprise-grade integrations and open standards ⁶
LangChain /LangGraph	Graph-Based Workflows & State Management	StateGraph, Nodes, Edges (Normal & Conditional)	Stateful, long-running applications that require fine-grained control over execution and dynamic decision-making within a structured loop ¹⁹	Persistence layer is supported out-of-the-box in the Platform version ²³	Yes, as the framework's blog highlights its role in defining context ¹⁹

Roo Code	Local-First, Autonomous IDE Agent	Modes (Code, Architect, Debug), Tools, Task Todo List	Privacy-sensitive projects, proprietary codebases, or scenarios where transparency and local execution are paramount ²⁴	A 'Task Todo List' feature serves as a form of integrated memory ²⁶	Yes, as a key part of its interoperability strategy ⁶
----------	-----------------------------------	---	--	--	--

The choice of framework is not a matter of feature superiority but of selecting a philosophical architecture that aligns with the specific problem. For instance, a team building a predictable, deterministic system should consider a framework like LangGraph for its explicit control over workflows.¹⁹ Conversely, a team building a system that requires a high degree of collaborative autonomy might find a framework like CrewAI more suitable for its role-based design.²⁰ A mature agentic system will likely be a hybrid, combining structured workflows for predictable tasks with autonomous agents for dynamic problem-solving.¹⁹ The ultimate best practice is to design a hybrid system that leverages the strengths of both approaches.

6. The Roo Code Framework: A Focused Case Study

6.1. Architectural Philosophy and Key Features

Roo Code represents a distinct architectural philosophy within the agentic ecosystem: a local-first, open-source AI agent that operates directly within the developer's editor.¹⁶ This is a direct counter-positioning to the cloud-based, black-box SaaS tools that dominate the market, with a strong emphasis on privacy and transparency.²⁵

Its key features include:

- **Local and Open-Source:** The agent runs locally, giving developers full control over what

their AI sees and sends.²⁵ This makes it an ideal tool for projects that are private, confidential, or subject to non-disclosure agreements.

- **Editor Integration:** Roo Code is designed to live inside the developer's codebase and integrate directly into editors like VS Code.¹⁶ It can read and write files, execute terminal commands, and perform multi-step tasks within the familiar development environment.¹⁶
- **Task Memory and Modes:** It includes features like a "Task Todo List" that acts as a simple, integrated task memory system for managing multi-step workflows.²⁶ It also supports different "modes" (e.g., Code, Architect, Debug) to tailor its behavior to specific tasks.¹⁶
- **Extensibility:** Roo Code supports multiple LLM providers, including local options like Ollama for completely offline operation.¹⁶ Its behavior can be further customized with .roorules files and APIs.²⁴

6.2. Roo Code and the Best Practices

Roo Code's design aligns closely with several emerging best practices. Its explicit support for the AGENTS.md standard is a significant finding.⁶ This allows a developer to use the broader, open standard for guiding their agents while benefiting from Roo Code's local, privacy-focused approach. This demonstrates how a simple, open standard can become the key enabler for a trend toward specialized, interoperable tools. Without

AGENTS.md, a team would be locked into a proprietary configuration format and would not be able to switch between tools with different value propositions. The convergence on this open standard prevents vendor lock-in and fosters a diverse ecosystem.⁶

Roo Code's local-first architecture is a direct response to a major, enterprise-level need for privacy and control over proprietary code. The fact that it is a popular and viable alternative to cloud-based solutions proves that a large segment of the market places a high value on transparency and hackability over polished user experience or speed.²⁵ This positions Roo Code not just as another tool but as a key indicator of a market trend driven by a need for security and self-hosting.

7. Conclusion: A Blueprint for Documentation and Implementation

The transition to agentic coding is not merely an incremental improvement in development tools; it is a fundamental shift in the software engineering discipline. This analysis of the latest best practices provides a clear blueprint for designing and documenting agentic systems that are robust, scalable, and maintainable.

The synthesis of this analysis yields several core recommendations for documentation and implementation:

- **Standardize with AGENTS.md:** The AGENTS.md standard is the single most important tool for ensuring consistency and interoperability in an agent-driven development environment. Its adoption as an open, machine-readable guide is critical for accelerating the onboarding of AI agents and maintaining high standards for code quality across a team.⁵ The detailed sections on its structure and implementation will form the basis of a documentation set's "How-To Guides," providing a clear, actionable path for teams to follow.
- **Design for Stateful Memory:** To move beyond basic, single-session interactions, all agentic projects should explicitly plan for long-term knowledge retention. This involves adopting a tiered memory architecture that combines ephemeral session history with persistent, structured knowledge bases.¹² The detailed explanations of memory architectures and the A-MEM framework's novel, agent-driven approach can serve as the "Explanation" domain of a documentation set, providing the necessary theoretical knowledge for developers to build truly autonomous systems.¹⁸
- **Embrace Architectural Nuance:** The choice of an agentic framework is a strategic decision that depends on the problem domain. The report's analysis of the "agents vs. workflows" dichotomy and its comparison of leading frameworks provides a decision-making framework.¹⁹ This information can be used to create the "Reference" domain of a documentation set, such as a scannable table, that helps developers quickly map a framework to a specific use case.¹⁹

By adopting these best practices, a development team can create a Diátaxis documentation set that is not just a collection of technical facts but a comprehensive, principled guide for navigating the agentic revolution. A concluding tutorial that walks a user through a full agentic workflow, such as the bug triage example, can tie all the concepts together, providing a practical, goal-oriented learning experience.¹ The future of software development lies not in writing code, but in orchestrating intelligent, self-evolving systems.

Works cited

1. Agentic Programming: How to Stay Relevant as a Developer in an AI Future, accessed September 25, 2025, <https://jaystechbites.com/posts/2025/agentic-programming-developer-relevance-ai-future/>
2. Introducing an agentic coding experience in your IDE - AWS, accessed September 25, 2025,

- <https://aws.amazon.com/blogs/devops/amazon-q-developer-agentic-coding-experience/>
3. aws.amazon.com, accessed September 25, 2025, <https://aws.amazon.com/blogs/devops/amazon-q-developer-agentic-coding-experience/#:~:text=Agentic%20coding%20in%20the%20IDE.time%20through%20natural%20language%20conversations.>
 4. The Agentic AI Handbook: A Beginner's Guide to Autonomous Intelligent Agents, accessed September 25, 2025, <https://www.freecodecamp.org/news/the-agentic-ai-handbook/>
 5. Agents.md: The README for Your AI Coding Agents - Research AIMultiple, accessed September 25, 2025, <https://research.aimultiple.com/agents-md/>
 6. AGENTS.md Emerges as Open Standard for AI Coding Agents - InfoQ, accessed September 25, 2025, <https://www.infoq.com/news/2025/08/agents-md/>
 7. AGENTS.md: A New Standard for Unified Coding Agent Instructions - Addo Zhang - Medium, accessed September 25, 2025, <https://addozhang.medium.com/agents-md-a-new-standard-for-unified-coding-agent-instructions-0635fc5cb759>
 8. Improve your AI code output with AGENTS.md (+ my best tips) - Builder.io, accessed September 25, 2025, <https://www.builder.io/blog/agents-md>
 9. Introduction to Agents.md | genai-research - Weights & Biases - Wandb, accessed September 25, 2025, https://wandb.ai/wandb_fc/genai-research/reports/Introduction-to-Agents-md--VmlldzoxNDEwNDI2Ng
 10. AGENTS.md Explained: One File to Rule All Agents - YouTube, accessed September 25, 2025, <https://www.youtube.com/watch?v=TC7dK0gwgq0>
 11. Agents.md: A Comprehensive Guide to Agentic AI Collaboration | by DhanushKumar | Sep, 2025 | Artificial Intelligence in Plain English, accessed September 25, 2025, <https://ai.plainenglish.io/agents-md-a-comprehensive-guide-to-agentic-ai-collaboration-571df0e78ccc>
 12. Stateful Agents: The Missing Link in LLM Intelligence | Letta, accessed September 25, 2025, <https://www.letta.com/blog/stateful-agents>
 13. What is long context and why does it matter for AI? | Google Cloud Blog, accessed September 25, 2025, <https://cloud.google.com/transform/the-prompt-what-are-long-context-windows-and-why-do-they-matter>
 14. Memory and State in LLM Applications - Arize AI, accessed September 25, 2025, <https://arize.com/blog/memory-and-state-in-llm-applications/>
 15. [Discussion] AGENT.md is only half the stack. Where's the plan for project memory? - Reddit, accessed September 25, 2025, https://www.reddit.com/r/ClaudeCode/comments/1mo8w3b/discussion_agentmd_is_only_half_the_stack_wheres/
 16. Roo Code: A Guide With Seven Practical Examples - DataCamp, accessed September 25, 2025, <https://www.datacamp.com/tutorial/roo-code>
 17. Agentic AI frameworks for enterprise scale: A 2025 guide - Akka, accessed

- September 25, 2025, <https://akka.io/blog/agentic-ai-frameworks>
18. A-MEM: Agentic Memory for LLM Agents - arXiv, accessed September 25, 2025, <https://arxiv.org/abs/2502.12110>
 19. How to think about agent frameworks - LangChain Blog, accessed September 25, 2025, <https://blog.langchain.com/how-to-think-about-agent-frameworks/>
 20. AI Agent Frameworks: Choosing the Right Foundation for Your Business | IBM, accessed September 25, 2025, <https://www.ibm.com/think/insights/top-ai-agent-frameworks>
 21. 5 AI Agent Frameworks Compared - KDnuggets, accessed September 25, 2025, <https://www.kdnuggets.com/5-ai-agent-frameworks-compared>
 22. LangGraph - LangChain Blog, accessed September 25, 2025, <https://blog.langchain.com/langgraph/>
 23. LangGraph Platform - LangChain, accessed September 25, 2025, <https://www.langchain.com/langgraph-platform>
 24. Roo Code - AI Agent for Debugging and Code Completion, accessed September 25, 2025, <https://bestaiagents.ai/agent/roo-code>
 25. Roo Code: Open-Source AI That Works With Your Code, Not Against It - We Are Founders, accessed September 25, 2025, <https://www.wearefounders.uk/roo-code/>
 26. Roo Code Docs | Roo Code Documentation, accessed September 25, 2025, <https://docs.roocode.com/>
 27. Introduction - CrewAI, accessed September 25, 2025, <https://docs.crewai.com/introduction>